

# JDBC TIPS AND TRICKS

Montreal Meetup Jan 21/2020

Dave Cramer  
January 2020



# Crunchy Data

Leading provider of trusted open source PostgreSQL technology, support and training.



*Powering Innovation With The World's Most  
Advanced Open Source Database*

# Introduction

- Dave Cramer
- Maintainer for the JDBC driver since 1999
- There are many options for connecting
- Many of them I didn't totally understand
- This talk hopes to unveil some of the more interesting ones
- And explain how they work.

# Overview

- Connecting to the driver
- Connection options that change the behaviour
- Performance tips
- Logical Decoding

# Connecting to the server

# Connection Properties

- PG\_DBNAME
- PG\_DBHOST
- PG\_DBPORT
- USER
- PASSWORD

These can be used in the following manner

```
Properties props = new Properties();  
props.setProperty(PGProperty.PG_DBNAME.getName(), "test");  
props.setProperty(PGProperty.PG_HOST.getName(), "localhost");  
props.setProperty(PGProperty.PG_PORT.getName(), "5432");  
props.setProperty(PGProperty.USER.getName(), "davec");  
props.setProperty(PGProperty.PASSWORD.getName(), "");  
Connection connection = DriverManager.getConnection("jdbc:postgresql:", props);
```

# URL options

- jdbc:postgresql:
  - Connects to localhost, port 5432, database specified in user
- jdbc:postgresql://host/
  - Connects to <host>, port 5432, and database specified in user
- jdbc:postgresql://host:port/
  - Connects to <host><port> and database specified in user
- jdbc:postgresql:database
  - Connects to localhost, port 5432 database “database”
- jdbc:postgresql://host:port/database
  - Specify the host, port and database

# Connection fail over

- jdbc:postgresql://host1:port, host2:port/database?targetType=primary|secondary|preferSecondary
- Allows you to specify multiple servers for the driver to connect to
- Provides a way to prefer a server type.



# Connection Failover tuning

- `targetServerType=primary, secondary, preferSecondary`
- Observes if server allows writes to chose
- `loadBalanceHosts=boolean` will randomly pick from suitable candidates
- `hostRecheckSeconds=number of seconds between checking status (read or write) of hosts` default is 10 seconds

# Logging

- `loggerLevel = OFF | DEBUG | TRACE`
  - Enables `java.util.logging.Logger` `DEBUG=FINE`, `TRACE=FINEST`
  - Not intended for SQL logging but rather to debug the driver
- `loggerFile=<filename>` the file to output the log to. If this is not set then the output will be written to the console.

# Logging

- `FINEST: FE=> SimpleQuery(query="select 1")`
- `Mar 09, 2018 1:20:33 PM org.postgresql.core.v3.QueryExecutorImpl receiveFields`
- `FINEST: <=BE RowDescription(1)`
- `Mar 09, 2018 1:20:33 PM org.postgresql.core.v3.QueryExecutorImpl receiveFields`
- `FINEST: Field(?column?,INT4,4,T)`
- `Mar 09, 2018 1:20:33 PM org.postgresql.core.v3.QueryExecutorImpl processResults`
- `FINEST: <=BE DataRow(len=1)`

# Logging continued

- We will honour `DriverManager.setLogStream` or `DriverManager.setLogWriter`
- Parent logger is `org.postgresql`
- Since we are using `java.util.Logging`, we can use a properties file to configure logging

# Logging properties file

- `handlers=java.util.logging.FileHandler`
- `org.postgresql.level=FINEST`
- `java -Djava.util.logging.config.file=...`
- `handlers=java.util.logging.ConsoleHandler`
- `org.postgresql.level=ALL`
- `org.postgresql.Driver.level=INFO`
- `org.postgresql.core.v3.level=FINE`

# Find connection leaks

- `logUnclosedConnections=boolean`
- Provides an easy way to find connection leaks
- If this is turned on we track connection opening by saving a `Throwable` when the connection is opened. If the finalizer is reached and the connection is still open the stacktrace message created when the connection was opened is printed out.

# Autosave

- autosave = never | always | conservative
- PostgreSQL transaction semantics are all or nothing. This is not always desirable
- autosave=always will create a savepoint for every statement in a transaction.
- The effect of which means that if you do
  - Insert into invoice\_header ...
  - Insert into invoice\_lineitem ...
  - If the insert into invoice\_lineitem fails the header will still be valid.
- In conservative mode if the driver determines that reparsing the query will work then it will be reparsed and retried.

# Binary Transfer

- `binaryTransferEnable`=comma separated list of oid's or names
- `binaryTransferDisable`
- Currently the driver will use binary mode for most built-in types.



# SimpleQueryMode

- Client sends an SQL Command(s)
- Server replies with RowDescription
  - Each column has a name
  - Oid of table
  - Attribute (column) number
  - Type Oid
  - Binary or Text
- Server sends all of the data rows
- Server CommandComplete and ReadyForQuery

# ExtendedQuery

- Parse
  - Send query string with placeholders, and parameter types (can be named)
  - Response is ParseOK
- Bind
  - Each parameter has format(binary or text) and value
  - Response is BindOK

# ExtendedQuery

- Describe (RowDescription)
  - Response is RowDescriptions
  - Number of fields, field name, type, text or binary
- Execute
  - Responds with DataRows and CommandComplete
- Sync
  - Responds with ReadyForQuery

# preferQueryMode

- simple
  - Fewer round trips to db no bind, no parse
  - Required for replication connection
- extended
  - Default creates a server prepared statement, uses parse, bind and execute.
  - Protects against SQL injection
  - Possible to re-use the statement

# preferQueryMode

- extendedForPrepared
  - Does not use extended for statements, only prepared statements
  - Potentially faster execution of statements
- extendedCacheEverything
  - Uses extended and caches even simple statements such as 'select a from tbl' which is normally not cached

# defaultRowFetchSize=int

- Default is 0 which means fetch all rows
  - This is sometimes surprising and can result in out of memory errors
- If set `*AND* autocommit=false` THEN will limit the number of rows per fetch
- Potentially significant performance boost

# stringtype=varchar|unspecified

- Default is varchar, which tells the server that strings are actually strings
- You can use stringtype='unspecified'
  - Useful if you have an existing application that uses `setString('1234')` to set an integer column.
  - Server will attempt to cast the string to the appropriate type.

# ApplicationName=String

- sets the application name
- Servers version 9.0 and greater
- Useful for logging and seeing which connections are yours in pg\_stat\_activity, etc.



# readOnly=boolean

- The default is false
- True sends `SET SESSION CHARACTERISTICS AS TRANSACTION READ ONLY` to the server.
- This blocks any writes to persistent tables, interestingly you can still write to a temporary table.

# disableColumnSanitizer=boolean

- columnSanitizer folds column names to lower case.
- Column names like FirstName become firstname.
- Resultset.getInt("firstname")
- default is to sanitize names

# assumeMinServerVersion=String

- Currently there are only 2 use cases
- 9.0 which will enable
- ApplicationName=ApplicationName (defaults to PostgreSQL JDBC Driver)
- sets extra float digits to 3
- 9.4 necessary for replication connections

# currentSchema=someschema

- by default the current schema will be “public”
- If you want to refer to a table in a different schema it would have to be specified by schema.table
- If you set this connection property to “audit” for example instead of “select \* from audit.log” you could use select \* from log;

# reWriteBatchedInserts=true

- Enables the driver to optimize batch inserts by changing multiple insert statements into one insert statement.
- insert into tab1 values (1,2,3);
- insert into tab1 values (4,5,6);
- Rewritten as “insert into tab1 values (1,2,3), ..(4,5,6)

# replication=false, true, database

- True tells the backend to go into walsender mode
- Setting to database enables logical replication for that database
- Simple query mode, subset of commands
- Must be accompanied by `assumMinServerVersion="9.4"` and `preferQueryMode="simple"`

# Performance tricks

- `setFetchSize`
- `rewriteBatchInserts`

# Set FetchSize performance

- Fetch a large amount of data with different fetch sizes

```
public static final String QUERY = "SELECT t FROM number";
```

```
@Benchmark
```

```
public void test(Blackhole blackhole, PgStatStatements pgStatStatements) throws SQLException {  
    pgStatStatements.setTestName(QueryBenchmarks.JMHTestNameFromClass(_6_String_NoAutocommit.class));
```


```
    QueryUtil.executeProcessQueryNoAutocommit(QUERY, resultSet -> {  
        while (resultSet.next()) {  
            blackhole.consume(resultSet.getString(1));  
        }  
    });  
}
```

```
// Used to fetch rows in batches from the db. Will only work if the connection does not use  
AutoCommit
```

```
PGProperty.DEFAULT_ROW_FETCH_SIZE.set(properties, FETCH_SIZE);
```



# Time it takes to fetch 1M rows



# What are the options for inserting lots of data

- For each row insertExecute this is the slowest
- For each row insertBatch this would be ideal
- Insert into foo (i,j) values (1,'one'), (2,'two') .... (n,'n') hand rolled code
- Copy into foo from stdin...

# JDBC micro benchmark suite

- Java 1.8\_60
- Core i7 2.8GHz
- PostgreSQL 9.6
- <https://github.com/pgjdbc/pgjdbc/tree/master/ubenchmark>
- create table batch\_perf\_test(a int4, b varchar(100), c int4)

Table "public.batch\_perf\_test"

Column	Type
a	integer
b	character varying(100)
c	integer

# INSERT Batch 1 row at a time

- For each row Insert into perf (a,b,c) values (?, ?, ?)
- After N rows executeBatch
- Normal mode this executes N inserts, not any faster than
- Looping over N inserts without batch mode

# INSERT Batch N rows\_at\_a\_time

- For each row Insert into perf (a,b,c) values (?, ?, ?), (?, ?, ?), (?, ?, ?), (?, ?, ?)
- After N/ rows\_at\_a\_time rows executeBatch
- Given 1000 (N) rows if we insert them 100(rows\_at\_a\_time) , end up inserting 10 rows 100 wide
- More data inserted per statement, less statements

# INSERT Batch with insertRewrite

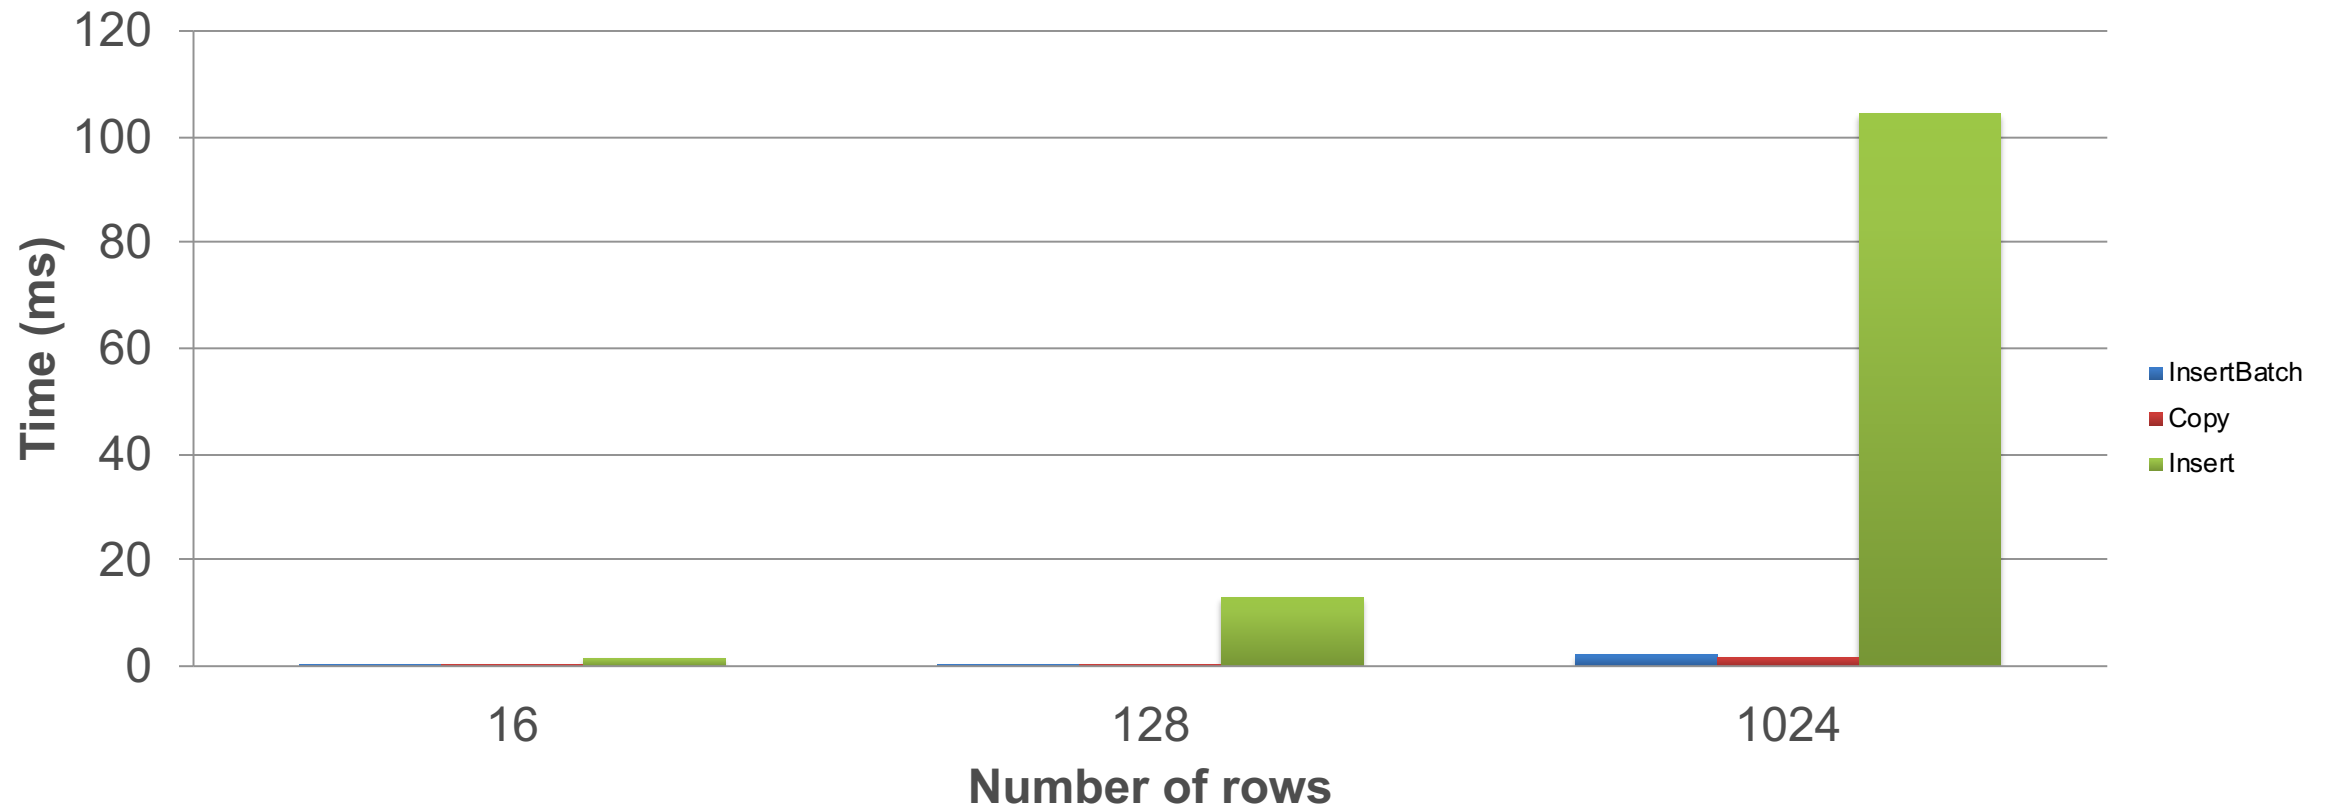
- For each row Insert into perf (a,b,c) values (?, ?, ?)
- After N rows executeBatch
- Same as last slide except we set the connection parameter insertRewrite=true
- As of version 1209 this has been enabled
- Same as insert into foo (i,j) values (1,'one'), (2,'two') .... (n,'n') except the driver does it for you.

# Copy

- Loop over the rows creating the input string in memory
- Build a string in memory which looks like  
0\t0\t0\n1\t1\t1\n....
- The string will end up being `nrows / rows_at_a_time` long
- Use the copy API to copy this into the table

# Results

Batch size of 128





# Conclusion

- Compared to batch inserts, plain inserts are very slow for large amounts of data

# How not to use JDBC (unfortunately typical)

- Open connection
- Prepare statement 'select \* from foo where id=?'
- `preparedStatment.executeQuery()`
- `preparedStatement.close()`
- Close Connection
- Without a pool connection creation is a heavyweight operation. PostgreSQL uses processes so each connection is a process
- Does not take advantage of caching

# Better solution

- Open connection
- Prepare statement 'select \* from foo where id=?'
- By default after 5 executions will create a named statement `PARSE S_1` as 'select \* from foo where id=?'
- Multiple `PreparedStatement.executeQuery()` `BIND/EXEC` instead of `PARSE/BIND/EXEC`
- Never close the statement if possible

# Query cache best practices

- Client side query cache only works in 9.4.1203 and up
- Do not use generated queries, as they generate new server side prepared statement
- Things like `executeUpdate('insert into foo (i,l,f,d) values (1,2,3,4)')` will never use a named statement
- Do not change the type of a parameter as this leads to DEALLOCATE/PREPARE
- `Pstmt.setInt(1,1)`
- `Pstmt.setNull(1,Types.VARCHAR)` this will cause the prepared statement to be deallocated

# Less obvious issues

- Server Prepare activated after 5 executions
- There is a configuration parameter called `prepareThreshold` (default 5)
- `PGStatement.isUseServerPrepare()` can be used to check
- After 5 executions of the same prepared statement we change from unnamed statements to named
- Named statements will use binary mode where possible;
- binary mode is faster when we have to parse things like timestamps
- Named statements are only parsed once on the server then bind/execute operations on the server

# setFetchSize

- If we don't use a fetch size we will read the entire response into memory then process
- Optimizing the data sent at one time reduces memory usage and GC
- Only works within a transaction
- Make sure fetch size is above 100
- If you have a lot of data this is really the only way to read it in without an Out Of Memory Exception

# Logical Replication Overview

- Reads the WAL logs and outputs them in any format you want
- Read changes
- Send confirmation of changes read
- GOTO read more changes

# Logical Replication High level Steps

- Create a replication connection
- Create a logical replication slot
- Read changes
- Send confirmation of changes read
- GOTO read more changes



# Create a Replication Connection

```
String url = "jdbc:postgresql://localhost:5432/postgres";  
Properties props = new Properties();  
PGProperty.USER.set(props, "postgres");  
PGProperty.PASSWORD.set(props, "postgres");  
PGProperty.ASSUME_MIN_SERVER_VERSION.set(props, "9.4");  
PGProperty.REPLICATION.set(props, "database");  
PGProperty.PREFER_QUERY_MODE.set(props, "simple");  
Connection con = DriverManager.getConnection(url, props);  
PGConnection replConnection=con.unwrap(PGConnection.class);
```

# Create a Logical Replication Slot

```
String outputPlugin = 'test_decode';
try (PreparedStatement preparedStatement =
    connection.prepareStatement("SELECT *
FROM pg_create_logical_replication_slot(?, ?)"))
{
    preparedStatement.setString(1, slotName);
    preparedStatement.setString(2, outputPlugin);
    preparedStatement.executeQuery()
}
```

# Create a replication stream

```
PGReplicationStream stream =  
    pgConnection  
        .getReplicationAPI()  
        .replicationStream()  
        .logical()  
        .withSlotName(SLOT_NAME)  
        .withStartPosition(lsn)  
        .withSlotOption("include-xids", true)  
        .start();
```

# Read Changes from database

```
while (true) {  
    //non blocking receive message  
    ByteBuffer msg = stream.readPending();  
    if (msg == null) {  
        TimeUnit.MILLISECONDS.sleep(10L);  
        continue;  
    }  
    int offset = msg.arrayOffset();  
    byte[] source = msg.array();  
    int length = source.length - offset;  
    System.out.println(new String(source, offset, length));  
    //feedback  
    stream.setAppliedLSN(stream.getLastReceiveLSN());  
    stream.setFlushedLSN(stream.getLastReceiveLSN());  
}
```

# <https://github.com/pgjdbc/pgjdbc>

- Credit where credit is due:
- Much of the optimization work on the driver was done by Vladimir Sitnikov
- Rewriting batch statements thanks to Jeremy Whiting
- Replication support was provided by Vladimir Gordiychuk
- Questions ?

# THANK YOU!

**Dave Cramer**  
**dave.cramer@crunchydata.ca**