



JDBC Performance from the Inside

Sept, 2016



Introduction

- Dave Cramer
- Work for Pivotal on the Greenplum database project <https://github.com/greenplum-db/gpdb>
- Maintainer for the JDBC driver since 1999
- Overheard conversation about how bad the driver was with respect to how it was built.
- Allegedly was very CPU intensive because of inheritance
- Things have changed a lot since then so I thought I'd show just how much better



Overview

- History of the driver
- Previous source layout
- Typical usage pattern
- Using Prepared Statements
- Batch processing how and why
- Optimal Fetch Size



History

- Originally written by Peter Mount in 1997
- Supported JDBC 1.2
- 1997 JDBC 1.2 Java 1.1
- 1999 JDBC 2.1 Java 1.2
- 2001 JDBC 3.0 Java 1.4
- 2006 JDBC 4.0 Java 6
- 2011 JDBC 4.1 Java 7
- 2014 JDBC 4.2 Java 8
- Each one of these were incremental additions to the interface
- Requiring additional concrete implementations of the spec to be implemented

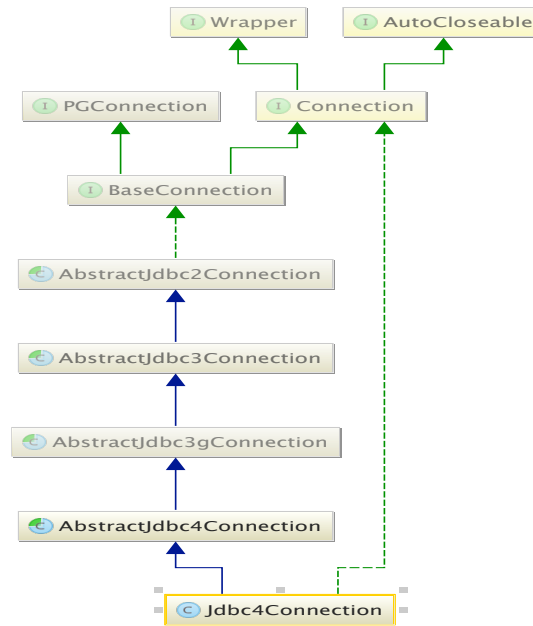
Source code layout



Before Maven (the ant years)

- jdbc2->jdbc3->jdbc3g->jdbc4->jdbc42
- Each one of these had abstract implementations, and concrete implementations
- Which one was built was determined by filters using ant
- Lions share of code was in jdbc2 package
- This meant that a concrete jdbc42 implementation extended AbstractJdbc42XXX which extended AbstractJDBC4XXX which extended AbstractJDBC3g ... all the way down to AbstractJDBC2XXX

Before Maven (the ant years)





Before Maven (the ant years)

- Why didn't you just use `-target` and compile previous versions with the latest compiler
- In theory since older versions of the spec will never attempt to access more recent interfaces this should "just work"
- Well embarrassingly we didn't think of it.
- Up until Java 8 this was possible.. The JDBC spec was never supposed to introduce a backward incompatibility.
- In Java 8 they added `java.time.*` the problem is: attempting to load a driver using an earlier JDK with `java.time` in it will cause a `ClassNotFoundException`. We are required to be able to pass a `java.time.*` object into `setObject`



Mavenizing the driver

- Why ? I was pretty hesitant to essentially rewrite the driver
- Easier to just include in your project

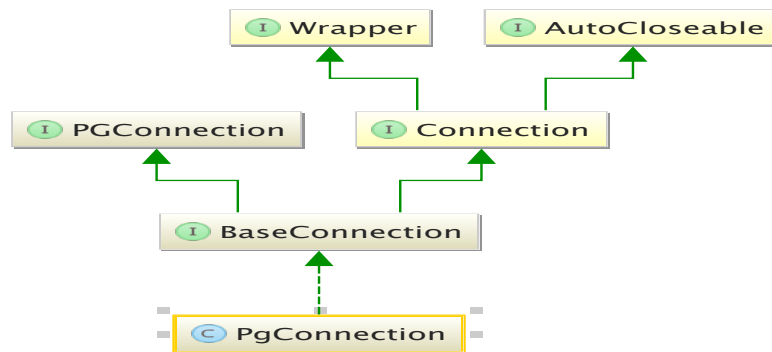
```
<dependency>  
  <groupId>org.postgresql</groupId>  
  <artifactId>postgresql</artifactId>  
  <version>9.4.1210</version> <!-- Java 8 →  
</dependency>
```
- Still have the same problems they are just solved differently
- Ant had filters to filter out which files are compiled for each build
- Maven uses pre-processing to add or remove code, avoids the multiple class extension
- *//#if mvn.project.property.postgresql.jdbc.spec >= "JDBC4.2"*



Mavenizing the driver

- Code re-organized release 1207 Dec 2015
- All of the abstract class machinations have been removed
- One class file works for all versions of JDBC
- Reduces CPU load
- Real advantage to mavenizing the project. The code is much simpler.
- Easier to debug, can be loaded into an IDE
- More people have provided Pull Requests

After Maven





Time to test the hypothesis

- Borrowed some code from <https://github.com/8kdata/javapgperf>
- Plug for ToroDB <https://github.com/torodb/torodb>
- **CREATE TABLE IF NOT EXISTS** number **AS**
SELECT i, 'Hello there ' || i **AS** t, '{"i": ' || i || ', "t": "' || 'Hello there ' || i || '" }'
AS j
FROM *generate_series*(1,10 * 1000 * 1000) **AS** i;

- Column | Type | Modifiers

```
-----+-----+-----  
i      | integer |  
t      | text    |  
j      | text    |
```



What do the tests do ?

- 1_int select i from number
- 2_String select t from number
- 3_IntString select i,t from number
- 4_IntStringJson select i, t, j from number
- 5_IntStringColumnNumber select i,t and use column number instead of column name
- 6_StringNoAutocommit select t from number
- Ran the test suite for a number of different JDBC versions 1200, 1202, 1204, 1208, 1210



What do the tests do ?

```
public class _4_IntStringJson {
    public static final String QUERY = "SELECT i, t, j FROM number";

    private class JsonElements {
        private int i;
        private String t;
    }

    @Benchmark
    public void test(Blackhole blackhole, PgStatStatements pgStatStatements) throws SQLException {
        pgStatStatements.setTestName(QueryBenchmarks.JMHTestNameFromClass(_4_IntStringJson.class));

        Gson gson = new Gson();

        QueryUtil.executeProcessQuery(QUERY, resultSet -> {
            while (resultSet.next()) {
                blackhole.consume(resultSet.getInt("i"));
                blackhole.consume(resultSet.getString("t"));
                blackhole.consume(gson.fromJson(resultSet.getString("j"), JsonElements.class));
            }
        });
    }
}
```



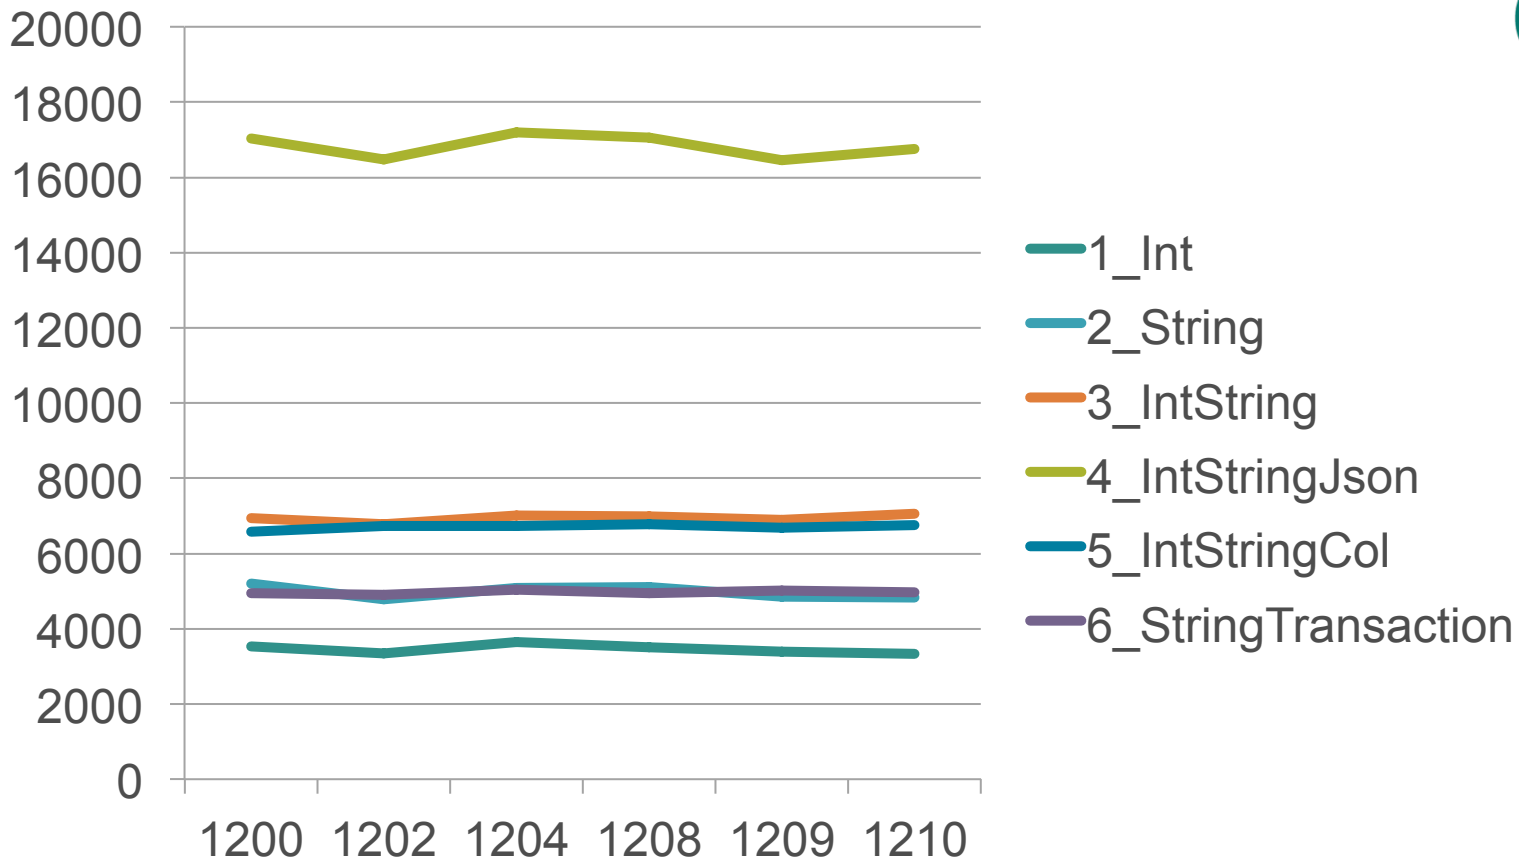
```
public class Main {
    private static final int ITERATIONS = 20;
    // Help profiling with sampling agents
    private static final int NO_FORKS_RUN_ON_THE_SAME_JVM = 0;

    public static void main(String[] args) throws RunnerException {
        if(args.length != 1 || args[0] == null || args[0].isEmpty()) {
            System.exit(1);
        }
        String testName = args[0];

        Options opt = addTestToOptionsBuilder(new OptionsBuilder(), testName)
            .addProfiler(org.postgresql.benchmarkprofilers.FlightRecorderProfiler.class)
            // .forks(1)
            // .jvmArgsPrepend("-Xmx128m")
            // We need to avoid warmup iterations as they however counts towards total Postgres time
            .warmupIterations(0)
            .measurementIterations(ITERATIONS)
            .timeUnit(TimeUnit.MILLISECONDS)
            .mode(Mode.SingleShotTime)
            .verbosity(VerboseMode.SILENT)
            .build();

        Collection<RunResult> runResults = new Runner(opt).run();

        runResults.stream().forEach(runResult ->
            System.out.printf(
                "Java:%s\t%.2f\n",
                runResult.getParams().getBenchmark(),
                runResult.getPrimaryResult().getScore()
            );
        );
    }
}
```





Something good came out of Maven

- Well it didn't really improve performance but it is:
- Easier to understand
- More people working on it
- Easier to work with simply import the maven pom.xml into IntelliJ
- Easier to push to maven



Some things that really did improve performance


- Set Fetch Size
- Fixed deadlocks
- Insert rewrite
- Fixing bugs



Set FetchSize performance

- Used exactly the same code as above with different fetch sizes

setFetchSize





Inserting batch Deadlock

- Ideally we would like to:
- Insert N rows of inserts where N is some arbitrarily large number
- PARSE S_1
- BIND/EXEC N TIMES
- DEALLOCATE



Unfortunately this doesn't work

- Driver is busy sending data, so it hasn't retrieved any responses
- Server is busy sending responses, so it can't fetch any more insert queries



Every so often we have to sync

- Parse S_1
- BIND/EXEC
- BIND/EXEC
- SYNC ... flush and wait for response
- The more sync's the slower it performs



Options for Inserting lots of data

- For each row `insertExecute`
- For each row `insertBatch`
- Insert values `(row1), (row2), ... (rowN)` hand rolled code
- `copy`



JDBC micro benchmark suite

- Java 1.8_60
- Core i7 2.8GHz
- PostgreSQL 9.6 (beta1)
- <https://github.com/pgjdbc/pgjdbc/tree/master/ubenchmark>
- create table batch_perf_test(a int4, b varchar(100), c int4)

Table "public.batch_perf_test"

Column	Type	Modifiers
--------	------	-----------

-----+-----+-----		
-------------------	--	--

a	integer	
b	character varying(100)	
c	integer	



The Code

- ```
public int[] insertBatch() throws SQLException {
 if (p2multi > 1) {
 // Multi values(),(),() case
 for (int i = 0; i < p1nrows;) {
 for (int k = 0, pos = 1; k < p2multi; k++, i++) {
 ps.setInt(pos, i);
 pos++;
 ps.setString(pos, strings[i]);
 pos++;
 ps.setInt(pos, i);
 pos++;
 }
 ps.addBatch();
 }
 }
}
```
- If we have 10 rows and p2multi is 2 the outer loop is executed 5 times and we insert 2 rows at a time
- Insert into foo (a,b,c) values (?, ?, ?), (?, ?, ?)



# INSERT Batch where p1multi =1

- For each row Insert into perf (a,b,c) values (?, ?, ?)
- After N rows executeBatch



# INSERT Batch where $p1multi > 1$

- For each row Insert into perf (a,b,c) values (?, ?, ?), (?, ?, ?), (?, ?, ?), (?, ?, ?)
- After  $N/p2multi$  rows executeBatch
- More data inserted per statement, less statements



# INSERT Batch with insertRewrite

- For each row Insert into perf (a,b,c) values (?, ?, ?)
- After N rows executeBatch
- Same as insertBatch except we set the connection parameter insertRewrite=true
- As of 1209 this is has been enabled
- Same as the previous slide except the driver does it for you.



# Copy

- Copy `batch_perf_test` from `stdin`  
1 s1 1  
2 s2 2  
3 s3 3
- Loop over the rows creating the input string in memory



# Hand rolled insert struct

- Insert into batch\_perf\_test select \* from unnest (?::batch\_perf\_test[])
- For N rows        setString to '{“(1,s1,1)”,“(2,s2,2)”,“(3,s3,3)”}'
- Add Batch
- executeBatch
- The query that gets executes look like:

Insert into batch\_perf\_test


select \*

from unnest ('{“(1,s1,1)”,“(2,s2,2)”,“(3,s3,3)”}'::batch\_perf\_test[])

# Results



## Batch size of 128








# Conclusion

- Compared to batch inserts, plain inserts are very slow for large amounts of data

# Results



## 1024 rows different batch sizes





# Bug fixes

- <https://github.com/pgjdbc/pgjdbc/pull/380>
- `QUERY_FORCE_DESCRIBE_PORTAL` shared the same value as `QUERY_DISABLE_BATCHING` effectively disabling batch inserts
- 10x increase in throughput

# How not to use JDBC (unfortunately typical)



- Open connection
- Prepare statement 'select \* from foo where id=?'
- `preparedStatment.executeQuery()`
- `preparedStatement.close()`
- Close Connection
- Without a pool connection creation is a heavyweight operation. PostgreSQL uses processes so each connection is a process
- Does not take advantage of caching



# Better solution

- Open connection
- Prepare statement 'select \* from foo where id=?'
- By default after 5 executions will create a named statement PARSE S\_1 as 'select \* from foo where id=?'
- Multiple preparedStatment.executeQuery() BIND/EXEC instead of PARSE/BIND/EXEC
- Never close the statement if possible



# Query cache best practices

- Client side query cache only works in 9.4.1203 and up
- Do not use generated queries, as they generate new server side prepared statement
- Things like `executeUpdate('insert into foo (i,l,f,d) values (1,2,3,4)')` will never use a named statement
- Do not change the type of a parameter as this leads to DEALLOCATE/PREPARE
- `Pstmt.setInt(1,1)`
- `Pstmt.setNull(1,Types.VARCHAR)` this will cause the prepared statement to be deallocated



# Less obvious issues

- Server Prepare activated after 5 executions
- There is a configuration parameter called `prepareThreshold` (default 5)
- `PGStatement.isUseServerPrepare()` can be used to check
- After 5 executions of the same prepared statement we change from unnamed statements to named
- Named statements will use binary mode where possible;
- binary mode is faster when we have to parse things like timestamps
- Named statements are only parsed once on the server then bind/execute operations on the server



# setFetchSize

- If we don't use a fetch size we will read the entire response into memory then process
- Performance numbers without fetch size 4825ms with 3686ms (fetch size set to 1000), 3161ms (fetch size 10000)
- Reducing the data sent at one time reduces memory usage and GC
- Only works with in a transaction
- Make sure fetch size is above 100
- If you have a lot of data this is really the only way to read it in without an Out Of Memory Exception





# Performance enhancements review

- Cache parsed statements across PreparedStatement calls now don't have to parse the statement in java each time
- Execute Batch changed to not execute statement by statement bug in code disabled batching
- Rewrite Batched inserts rewrites inserts from multiple insert into foo (a,b,c) values (1,2,3) to insert into foo (a,b,c) values (1,2,3), (4,5,6) this provides 2x-3x speed up
- Avoid Calendar cloning provides 4x speed increase for setTimestamp pr 376



# Conclusions

- Using insert rewrite gives us a 2-3x performance increase for batch inserts
- Makes sense as it is one trip
- Use `setFetchSize(100)` or greater and use transactions
- Don't close prepared statements.



# <https://github.com/pgjdbc/pgjdbc>

- Credit where credit is due:
- Much of the optimization work on the driver was done by Vladimir Sitnikov
- Much (if not all ) of the work to convert the build to Maven was done by Stephen Nelson
- Rewriting batch statements thanks to Jeremy Whiting
- Questions ?