

Logical Decoding

Presented to:

Dave Cramer
October 2018



Crunchy Data

Leading provider of trusted open source PostgreSQL technology, support and training.



Powering Innovation With The World's Most Advanced Open Source Database

TL;DR Logical Decoding

- Change data capture built in
- Insert, update, delete captured
- Changes are streamed *after* the transaction commits
- Third party change data capture (debezium)

Overview

- Uses Write Ahead Log mechanism to capture the changes
- Each row that is changed is sent to the output plugin
- Unlogged and temp tables do not go through WAL so not decoded

Overview

- DDL is not decoded
- Must have a user with replication permissions
- Each replication stream corresponds to a replication slot. Slots provide fine grain control over each replication stream.
- Potential security issues. Replication user sees all changes

Plugins

- The API provides the data in a binary format
- The output plugin decodes the binary into a format usable by external services
- Output plugins are written in C as shared libraries

Plugins

- Default plugin is test_decoding
- JSON decoder Wal2json
<https://github.com/eulerto/wal2json>
- Protobuf decoder <https://github.com/xstevens/decoderbufs>

Change Data Capture

- Typically done by some combination of triggers, log tables and polling
- Triggers on tables insert changes into a “log table”
- Log table is polled for new data
- Causes bloat in the database, lots of inserts and deletes

Change Data Capture

- Changes are pushed by the server.
- Server does not delete the change until it is confirmed
- Can do logical replication of a single database, and or table.

Server requirements

- Configuration
 - `wal_level = logical` (at least logical)
 - `max_replication_slots` = some number greater than 0
 - `synchronous_commit = on` (if you want synchronous replication)
- have to have a replication user, and configure `pg_hba.conf`

SQL functions

- `pg_create_logical_replication_slot('slot_name', 'decoder_name', [temporary boolean])`
 - Creates a slot and associates the slot with the decoder
 - As of pg10 can be temporary
- `pg_logical_slot_get_changes('slot_name', null, null, null)`
 - Get any changes from the last request for changes

SQL functions

- There are get and peek functions
- Also ways to control the output format
 - Include-timestamp
 - Include-xid
- Control the number of changes, with upto_lsn, and n_entries

Care and feeding of slots

- It's important to drop the slot when finished with it
 - `pg_drop_replication_slot('slot_name')`
- The server will keep all the WAL associated with that slot unless the WAL are consumed and confirmed
- Can fill up your server and cause postgresql to stop.

Monitoring

- `pg_replication_slots`
 - Slot name, plugin, type, database
 - Active, active pid
 - Xmin oldest tx the db needs to retain (can't be vacuumed)
 - Restart and confirmed LSN(Log Sequence Number) or position in the WAL.

Monitoring

- pg_stat_replication

<https://www.postgresql.org/docs/10/static/monitoring-stats.html#PG-STAT-REPLICATION-VIEW>

- State of replication
- Lag write, flush, replay
- Location various LSN's sent, write, flush, replay
- Sync state async, sync, quorum
- Write, flush and replay
 - Written but not flushed to disk or applied
 - Written and flushed but not applied

Replica Identity

- Determines how much data is written to the WAL on UPDATE or DELETE
 - DEFAULT: old data are only written to the WAL when the PK is changed
 - FULL: old data are always written
 - NOTHING: old data is never written
 - USING INDEX: old data is written if the column in the index is updated
 - Index needs to be unique, and cannot contain null values, or expressions

JDBC support

- Logical replication is built in to the JDBC driver
- Same requirements. Replication user, wal_level, etc

Steps to setup

- Create a slot, this just uses the `pg_create_logical_replication_slot` function
- Open a replication connection.
 - Must be 9.4
 - Simple query mode
 - Replication mode is 'database' (this means logical replication)
- Get a replication stream
- Read the changes
- Demo

THANK YOU!

Dave Cramer
dave.cramer@crunchydata.ca