

Active-Active Replication

Presented to: Scale 17x

Dave Cramer
March 2019



Dave Cramer

- Senior Data Architect @ Crunchy Data
- Major Contributor to PostgreSQL
- Working with PostgreSQL since 1999
- Maintainer of the PostgreSQL JDBC Driver
- Helping out with PL/R, some other small projects
- Real interest in where things like reactive Java and PostgreSQL are going

Crunchy Data

Leading provider of trusted open source PostgreSQL technology, support and training.



*Powering Innovation With The World's Most
Advanced Open Source Database*

Active-Active replication

- What is it?
- Why do we need/want it?
- Symmetric-ds
- Different use cases
- What it doesn't do
- The good and the bad ?

What is active-active ?

- Read/Write operations on more than one database instance
- Updates are propagated to all instances
- Single source of truth
- Sounds simple, right ?
- What happens if there are conflicts ?

Why do we want it

- Local performance
- Data close to the application

High Availability

- Two (or more) identically configured instances.
- Both accept writes
- Near zero downtime when failing over. No change to configuration of the new primary

Symmetric-ds

- Open source project by jumpmind.com
- <https://www.symmetricds.org/>
- <https://github.com/JumpMind/symmetric-ds>
- Replication toolkit

Symmetric-ds

- Written in Java
- Uses JDBC for database access
- Web app
- Asynchronous
- Unfortunately uses triggers and a log table
- Need to manage bloat in Postgres

Symmetric-ds

- Java and JDBC
- This allows the software to run on multiple platforms
- Enables multiple database vendors (relatively) easily

Symmetric-ds

- Asynchronous
- Synchronous replication over significant distances would slow the primary down
- Does introduce the possibility of data loss

Symmetric-ds

- Scales to thousands of databases
- Very configurable
 - Can configure which tables, columns or even rows
 - Can configured to push or pull data

Beyond Replication

- Zero(near) downtime upgrade of PostgreSQL
- Migration to PostgreSQL
- Cloud provider lock in

Upgrade on disk format

- Upgrading major versions of PostgreSQL can be painful
- Even pg_upgrade requires **some** downtime
- Change to heap checksums is not currently possible with pg_upgrade
- Admittedly this can be done with logical replication

Migration to PostgreSQL

- Migration from other database vendors to PostgreSQL
- Uses JDBC so works with a number of vendors
- Oracle, MSSQL, MariaDB, DB2 to mention a few

Cloud provider lock in

- Most do not allow binary replication
- Difficult to get ports other than 80 or 443 exposed
- Because Symmetric-ds is a web app it is relatively simple to replicate from one cloud to another

More than active-active

- Can be configured as single or bi-directional
- Can choose which tables to sync
- Has filters and transforms
- Can have one table(s) going up and another down

Central office with stores

- Each store has a local database for performance
- Each store can sell the item for a different price
 - For instance when they want to have a sale
- The resulting sale would be replicated to the head office in order to aggregate corporate sales
- When new items become available central office can push new items to the stores

Features

- How to deal with some of the details of replication
 - Conflict resolution
 - Transformations
 - Filters
 - Monitoring

Conflict Detection

- Various conflict resolution schemes
- Primary Key – will not allow a row to be inserted
- Changed Data – primary key plus only data changed
- Old data – All the existing row data will be used

Conflict Detection

- Timestamp – If the Primary key plus the timestamp column are equal
- Version column – If the Primary key plus the version are equal

Conflict Resolution

- Manual
 - Fix the data and mark it resolved
- Ignore
 - Ignore the batch or the row depending on config
- Newer wins
 - Either the source or destination with the latest change, based on either a timestamp or version column

Transformations

- Data moving from the source to the destination can be transformed as it moves
- This can be done when it is read from the source
- Or when written to the destination

Transformations

- The simplest and default is just to copy the data to the destination column
- Possible to change columns
 - Id -> account_id for instance
- Possible to remove NULL's
 - NULL timestamp to now()

Transformations

- A column can be omitted
- Set a column to a constant value
- Set to a system variable
 - Date, timestamp, node_id (source or target), null,
- Previous value before updating
 - Auditing of changes ?

Transformations

- Additive transform. $\text{target} = \text{target} + \text{multiplier} (\text{source_new} - \text{source_old})$
- Substring – expression is 0,5 which is passed to the java substring function
- Left – copies left most characters
- Bleft – copies left most bytes
- Many more

Transformations

- Bean Shell
- Executes the script in the expression
- Has access to column name, current value, old value, channel id, source node, target node and can execute SQL.

Filters

- Applied when loading the data
- Can be used to determine if the new data will be replicated
- For instance do not replicate new prices that are less than the old price

Monitoring

- Machine resources: CPU, Memory, disk space
- Replication status
 - Batch Error
 - Batch Unsent – data ready to be sent
 - Data unrouted – data ready to be batched
- Nodes offline

Monitoring

- Log file – Can be used with log aggregators
- Send emails when an error occurs

Panacea ... not

- Does not replicate sequences
- This creates a problem if the sequence exists on both primaries
- So what happens: insert on one instance increments the sequence on that instance, subsequent insert on replica will fail with primary key failure (data will be inserted without using sequence)

Panacea ... not

- Solution to this is to increment sequences by 2 or more
- If you have 3 primaries you would have to have increment each by 3 and so on

Panacea ... not

- Does not replicate DDL
- API allows for moving DDL over as well as data
- It is possible to lose data as this is asynchronous

In Summary The Good

- Java, runs anywhere
- Cross database
- Web app, easy egress and ingress
- Cloud migration, avoid lock in
- REST API

The Bad

- It uses triggers and they (symmetric-ds) don't deal with bloat well. This is specific to PostgreSQL
- It's,a bit complicated to setup
- Jumpmind has a paid (pro) version with a GUI
- The REST API is not complete enough to build a GUI
- GPL

How To

- <https://info.crunchydata.com/blog/a-guide-to-building-an-active-active-postgresql-cluster>

THANK YOU!

Dave Cramer
dave.cramer@crunchydata.ca